

```

In [12]: %matplotlib inline
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import torch.nn as nn
import torch
import requests
from io import BytesIO
import tqdm

bucky_urls = [
    "https://cdn.vox-cdn.com/thumbor/3yCbVJHk2QeCNc47W6yaR2ok-0E=/0x0:4298x2
    "https://www.ncaa.com/_flysystem/public-s3/styles/large_16x9/public-s3/i
    "https://chancellor.wisc.edu/content/uploads/2018/08/Bucky_Parade_unveil
    "https://upload.wikimedia.org/wikipedia/commons/6/68/BuckinghamUBadger.j
]

def get_image_from_url(url):
    response = requests.get(url, headers={'User-Agent': 'Mozilla/5.0'})
    img = Image.open(BytesIO(response.content))
    img = img.resize((128,128))
    return img

def prepare_image_for_torch(img):
    img_arr = np.asarray(img) / 255.
    img_arr = torch.from_numpy(img_arr).float()
    return img_arr

```

```

In [2]: class AutoEncoder(nn.Module):
    def __init__(self, in_chans, hidden_chans=10):
        super().__init__()
        self.encoder = nn.Linear(in_chans, hidden_chans, bias=False)
        self.decoder = nn.Linear(hidden_chans, in_chans, bias=False)

        for m in self.modules():
            if isinstance(m, nn.Linear):
                nn.init.normal_(m.weight, mean=0, std=0.001)

    def forward(self, x):
        B,H,W,D = x.shape
        x = x.flatten(1)
        z = self.encoder(x)
        xpred = self.decoder(z)
        return xpred.view(B,H,W,D)

```

```

In [3]: def loss(pred, y):
    """MSE loss"""
    return (pred - y).pow(2).mean()

def fit(model, loss, images, iterations=300):
    """
    Helper function to train a model.

```

```

- images: NxWxHx3 tensor
- iterations: number of training iterations
"""

opt = torch.optim.Adam(model.parameters(), lr=0.0005)
for it in tqdm.tqdm(range(iterations)):
    pred = model(images)
    l = loss(pred, images)

    l.backward()
    opt.step()
    opt.zero_grad()
return model

@torch.no_grad()
def display_reconstructions(model):
    """
    Helper function to visualize reconstructions obtained through the autoencoder
    """
    for i, url in enumerate(bucky_urls):
        img = get_image_from_url(url)
        img_arr = prepare_image_for_torch(img)
        img_rec = model(img_arr[None])[0]

        plt.subplot(len(bucky_urls), 2, 1+i*2)
        plt.imshow(img)
        plt.axis(False)
        plt.subplot(len(bucky_urls), 2, 2+i*2)
        plt.imshow(img_rec.numpy().clip(0, 1))
        plt.axis(False)

```

### Question 2-a

```

In [4]: # Declare the model
autoencoder = AutoEncoder(128*128*3, 10)



# Prepare training data
img = get_image_from_url(bucky_urls[0])
img_arr = prepare_image_for_torch(img)[None]

# Fit the model
fit(autoencoder, loss, img_arr)

# Display results
display_reconstructions(autoencoder)

```

100%|██████████| 300/300 [00:00<00:00, 304.24it/s]





Findings: It seems to have transformed each image into the first one and changes the brightness depending on the source image

```
In [13]: # Part b)
# Step 1: Prepare all 3 images and stack them into a 3xHxWxD tensor
# Prepare training data
imgs = torch.stack([prepare_image_for_torch(get_image_from_url(url)) for url

# Step 2: Declare model, fit to data, visualize results
autoencoder = AutoEncoder(128*128*3, 10)
fit(autoencoder, loss, imgs)
display_reconstructions(autoencoder)
new_bucky_link = "https://en.wikipedia.org/wiki/File:BuckinghamUBadger.jpg"
```

100%|██████████| 300/300 [00:01<00:00, 217.16it/s]



Findings: All of the images are the same as the source image except the new image with appears to be a combination of all other images

(c) Why is the model not generalizing well to new images? Suggest two ideas to improve the generalization of the model.

#### Answer

We don't have enough data to make the model generalize well to new images.

Ways to improve:

1. Increase the Dataset Size & Diversity
2. Increase models hidden channels/hidden layers

```
In [20]: # Part d) and e)
# Step 1: Prepare all 3 images and stack them into a 3xHxWxD tensor
# Step 2: Rewrite fit_with_noise() so that, at each iteration, images are fi
# Step 3: Declare the model, fit to data, visualize results

def add_noise(img):
    img = img.clone()
    rows = img.shape[0]
    cols = img.shape[1]
    s = int(min(rows,cols)/20) # size of spot is 1/20 of smallest dimension

    for i in range(100):
        x = np.random.randint(cols-s)
        y = np.random.randint(rows-s)
```

```




        img[y:(y+s),x:(x+s)] = 0

    return img

for i, url in enumerate(bucky_urls[0:3]):
    img = get_image_from_url(url)
    img = prepare_image_for_torch(img)

    plt.subplot(1, 3, 1+i)
    plt.imshow(add_noise(img))
    plt.axis(False)

```



```

In [23]: @torch.no_grad()
def display_reconstructions_with_noise(model):
    for i, url in enumerate(bucky_urls):
        img = get_image_from_url(url)
        img_arr = prepare_image_for_torch(img)
        img_noise = add_noise(img_arr)
        img_rec = model(img_noise[None])[0]

        plt.subplot(len(bucky_urls),3,1+i*3)
        plt.imshow(img)
        plt.axis(False)
        plt.subplot(len(bucky_urls),3,2+i*3)
        plt.imshow(img_noise.numpy().clip(0, 1))
        plt.axis(False)
        plt.subplot(len(bucky_urls),3,3+i*3)
        plt.imshow(img_rec.numpy().clip(0, 1))
        plt.axis(False)

def fit_with_noise(model, loss_fn, images, optimizer, epochs=300):
    model.train()
    for epoch in range(epochs):
        noisy_imgs = torch.stack([add_noise(img) for img in images]) # Appl
        optimizer.zero_grad()
        recon_imgs = model(noisy_imgs) # Forward pass
        loss = loss_fn(recon_imgs, images) # Compare with clean images
        loss.backward()
        optimizer.step()

    if epoch % 50 == 0:
        print(f"Epoch {epoch}/{epochs}, Loss: {loss.item():.6f}")

```

```

autoencoder = AutoEncoder(128 * 128 * 3, 10) # Ensure correct model architecture
optimizer = torch.optim.Adam(autoencoder.parameters(), lr=0.001)
loss_fn = torch.nn.MSELoss()

# Reshape images into a flattened batch (if needed)
imgs = torch.stack([prepare_image_for_torch(get_image_from_url(url)) for url in urls])

fit_with_noise(autoencoder, loss_fn, imgs, optimizer, epochs=300)

display_reconstructions_with_noise(autoencoder)

```

Epoch 0/300, Loss: 0.266351  
 Epoch 50/300, Loss: 0.001055  
 Epoch 100/300, Loss: 0.001988  
 Epoch 150/300, Loss: 0.001066  
 Epoch 200/300, Loss: 0.000649  
 Epoch 250/300, Loss: 0.000541



In [ ]: