

problems: 2 and 3

## Question 2

Task \ Augmentation	A. Horizontal Flip	B. Random Rotation	C. Color Jitter	D. Gaussian Noise	E. Random Crop & Resize
1. Determining Image Orientation	x	✓	x	x	x
2. Arrow Direction Detection	x	✓	x	x	x
3. Image Cleanliness Classification	x	x	x	✓	x
4. Object Size Detection	x	x	x	x	✓
5. Handwritten Digit Recognition	x	✓	x	x	✓
6. Facial Recognition	✓	✓	✓	x	✓
7. Medical Tumor Detection	x	x	x	x	x
8. OCR for Printed Text	x	x	x	x	x
9. Animal Species Classification	✓	✓	✓	x	✓
10. Barcode Scanning	x	✓	x	x	✓

## Question 3

- A) Done
- B) The model did not converge after ten epochs
- C) The filters look like some of the images main colors
- D) It is WAYYYYYY faster to train on GPU

```
In [20]: import matplotlib.pyplot as plt
import numpy as np

import torch
import torch.nn as nn
```

Download and prepare the CIFAR10 dataset

The CIFAR10 dataset contains 60,000 color images in 10 classes, with 6,000 images in each class. The dataset is divided into 50,000 training images and 10,000 testing images. The classes are mutually exclusive and there is no overlap between them.

```
In [21]: from torchvision import datasets
from torchvision import transforms
transform = transforms.Compose([
    transforms.ToTensor(),
])
db_train = datasets.CIFAR10(root='./', train=True, transform=transform, download=True)
db_test = datasets.CIFAR10(root='./', train=False, transform=transform, download=True)
def subsample(db, n):
    idx = np.random.permutation(len(db))[:n]
    db.data = db.data[idx]
    db.targets = np.array(db.targets)[idx]
    return db

db_train = subsample(db_train, 5000)
db_test = subsample(db_test, 1000)

img, y = db_train[0]
print(img.shape, img.min(), img.max(), y)
```

Files already downloaded and verified

Files already downloaded and verified

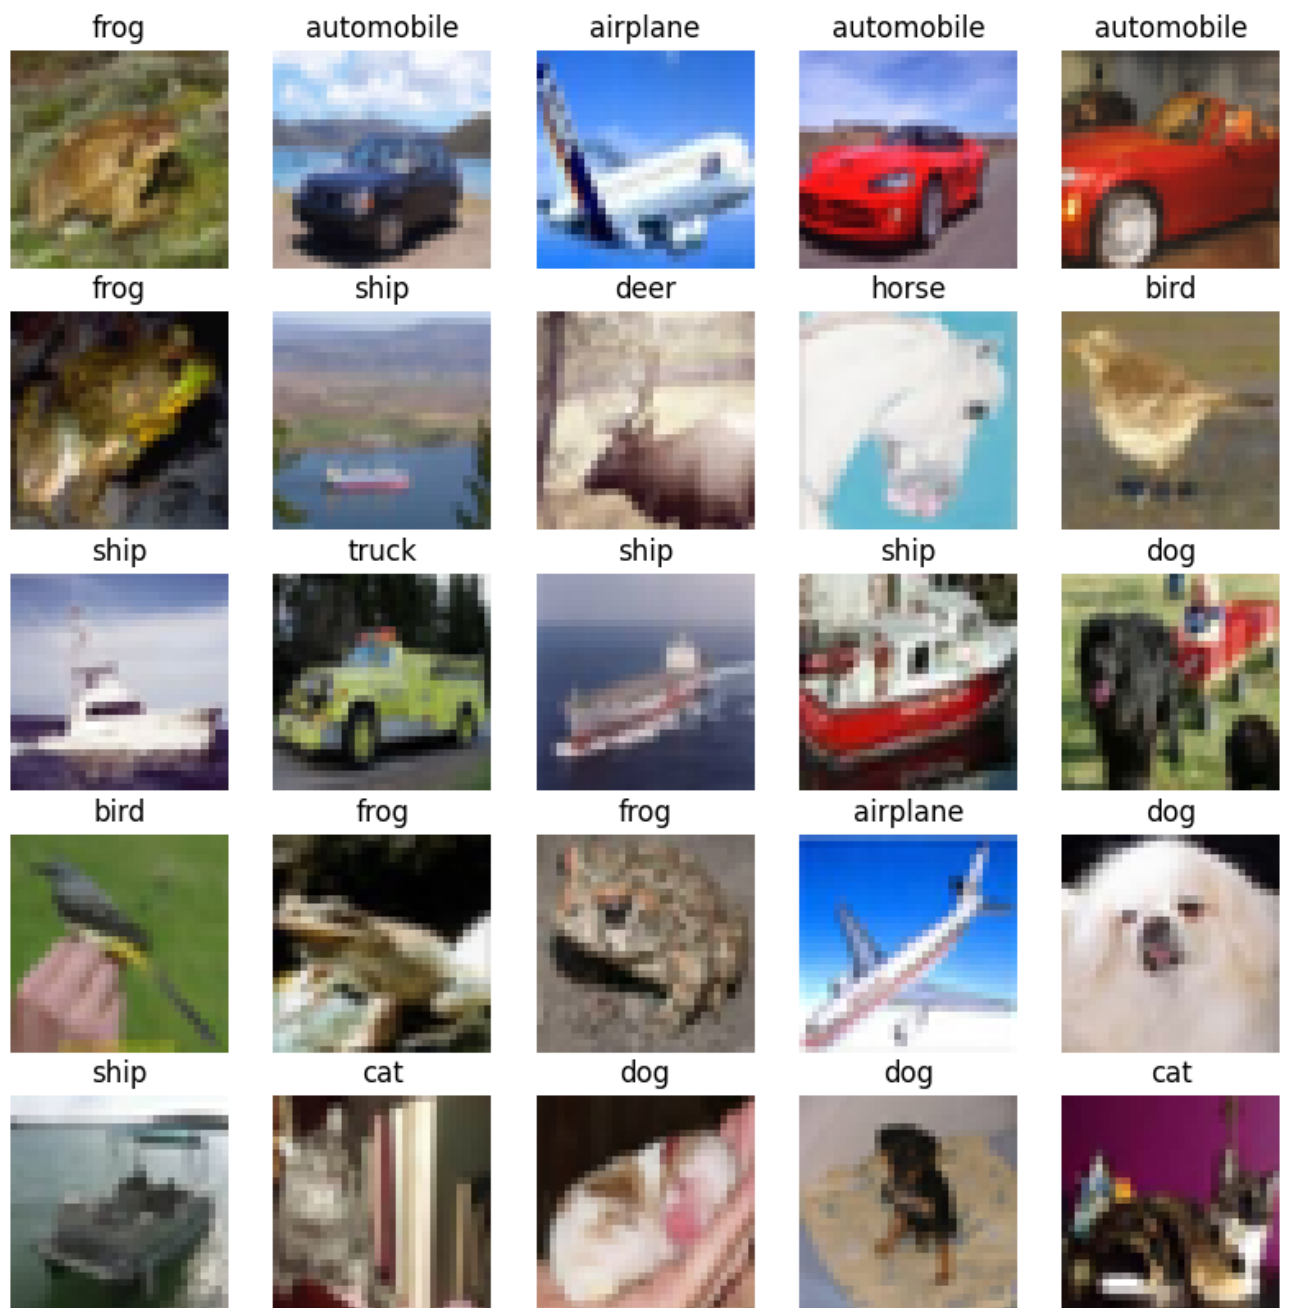
torch.Size([3, 32, 32]) tensor(0.0078) tensor(0.8157) 6

## Verify the data

To verify that the dataset looks correct, let's plot the first 25 images from the training set and display the class name below each image. Note that transforms.ToTensor transformed our image 32x32x3 image into a 3x32x32 tensor (the format that Conv2d layers expect). However, to display the image, we need to permute its dimensions again.

```
In [22]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                        'dog', 'frog', 'horse', 'ship', 'truck']

plt.figure(figsize=(10,10))
for i in range(25):
    img, lbl = db_train[i]
    plt.subplot(5,5,i+1)
    plt.imshow(img.permute(1, 2, 0))
    plt.title(class_names[lbl])
    plt.axis(False)
plt.show()
```



## Create the convolutional base

The 6 lines of code below define the convolutional base using a common pattern: a stack of `Conv2D` and `MaxPooling2D` layers.

As input, a CNN takes tensors of shape (image\_height, image\_width, color\_channels), ignoring the batch size. If you are new to these dimensions, `color_channels` refers to (R,G,B). In this example, you will configure your CNN to process inputs of shape (32, 32, 3), which is the format of CIFAR images. You can do this by passing the argument `input_shape` to your first layer.

```
In [23]: model = nn.Sequential(
          nn.Conv2d(3, 32, kernel_size=(5, 5)),
```

```

        nn.ReLU(),
        nn.MaxPool2d(kernel_size=(2, 2), stride=(2, 2)),
        nn.Conv2d(32, 64, kernel_size=(3, 3)),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=(2, 2), stride=(2, 2)),
        nn.Conv2d(64, 64, kernel_size=(3, 3)),
        nn.ReLU(),
    )
    dummy_img = torch.randn(1, 3, 32, 32)
    print(model(dummy_img).shape)

```

```
torch.Size([1, 64, 4, 4])
```

The output of every Conv2d and MaxPool2d layer is a 3D tensor of shape (channels, height, width). The width and height dimensions tend to shrink as you go deeper in the network. The number of output channels for each Conv2d layer is controlled by the second argument (e.g., 32 or 64). Typically, as the width and height shrink, you can afford (computationally) to add more output channels in each Conv2d layer. By feeding a dummy image, we observe that the model so far will output a (64, 4, 4) tensor.

## Add Dense layers on top

To complete the model, you will feed the last output tensor from the convolutional base of shape (64, 4, 4) into one or more Linear layers to perform classification. Linear layers take vectors as input (which are 1D), while the current output is a 3D tensor. First, you will flatten the 3D output to 1D, then add one or more Linear layers on top. CIFAR has 10 output classes, so you use a final Linear layer with 10 outputs. Since the convolutional base outputs 64x4x4 tensors, after flattening, we will have 1024 input features to the Linear layer.

```

In [24]: model.add_module("flatten", nn.Flatten())
         model.add_module("classifier", nn.Linear(1024, 10))

```

Here's the complete architecture of your model:

```

In [25]: print(model)

```

```

Sequential(
  (0): Conv2d(3, 32, kernel_size=(5, 5), stride=(1, 1))
  (1): ReLU()
  (2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
  (4): ReLU()
  (5): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1, ceil_mode=False)
  (6): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1))
  (7): ReLU()
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (classifier): Linear(in_features=1024, out_features=10, bias=True)
)

```

## Compile and train the model

```

In [26]: from torch.utils.data import DataLoader
import tqdm

def init_weights(model):
    for m in model.modules():
        if isinstance(m, (nn.Conv2d, nn.Linear)):
            nn.init.normal_(m.weight, mean=0, std=0.001)

def fit_one_epoch(model, opt, loader):
    losses, accuracies = [], []
    for images, labels in tqdm.tqdm(loader):
        pred = model(images)
        l = loss(pred, labels)
        acc = (pred.argmax(1) == labels).float().mean()

        l.backward()
        opt.step()
        opt.zero_grad()

        losses.append(l.detach().item())
        accuracies.append(acc.detach().item())
    return np.mean(losses), np.mean(accuracies)

@torch.no_grad()
def eval(model, loader):
    accuracies = []
    for images, labels in tqdm.tqdm(loader):
        pred = model(images)
        acc = (pred.argmax(1) == labels).float().mean()
        accuracies.append(acc.detach().item())
    return np.mean(accuracies)

def fit(model, loader_train, loader_test, epochs=50):
    opt = torch.optim.Adam(model.parameters(), lr=0.0005)
    hist_tr_loss, hist_tr_acc, hist_te_acc = [], [], []

```

```

for epoch in range(epochs):
    tr_l, tr_acc = fit_one_epoch(model, opt, loader_train)
    te_acc = eval(model, loader_test)

    print(f"Finished epoch {epoch} of {epochs}: Train Loss = {tr_l:.3f}")
    hist_tr_loss.append(tr_l)
    hist_tr_acc.append(tr_acc)
    hist_te_acc.append(te_acc)
    return hist_tr_loss, hist_tr_acc, hist_te_acc

loader_train = DataLoader(db_train, batch_size=64, shuffle=True, drop_last=True)
loader_test = DataLoader(db_test, batch_size=64, shuffle=False)
loss = nn.CrossEntropyLoss()
init_weights(model)
hist_tr_loss, hist_tr_acc, hist_te_acc = fit(model, loader_train, loader_test)

```

100%|██████████| 78/78 [00:03<00:00, 24.00it/s]

100%|██████████| 16/16 [00:00<00:00, 63.17it/s]

Finished epoch 0 of 50: Train Loss = 2.293    Train Acc = 0.109    Test Acc = 0.101

100%|██████████| 78/78 [00:03<00:00, 25.66it/s]

100%|██████████| 16/16 [00:00<00:00, 59.72it/s]

Finished epoch 1 of 50: Train Loss = 2.224    Train Acc = 0.135    Test Acc = 0.172

100%|██████████| 78/78 [00:02<00:00, 26.00it/s]

100%|██████████| 16/16 [00:00<00:00, 62.44it/s]

Finished epoch 2 of 50: Train Loss = 2.172    Train Acc = 0.199    Test Acc = 0.234

100%|██████████| 78/78 [00:03<00:00, 25.22it/s]

100%|██████████| 16/16 [00:00<00:00, 63.71it/s]

Finished epoch 3 of 50: Train Loss = 2.050    Train Acc = 0.224    Test Acc = 0.237

100%|██████████| 78/78 [00:03<00:00, 25.37it/s]

100%|██████████| 16/16 [00:00<00:00, 58.16it/s]

Finished epoch 4 of 50: Train Loss = 2.009    Train Acc = 0.242    Test Acc = 0.247

100%|██████████| 78/78 [00:03<00:00, 23.79it/s]

100%|██████████| 16/16 [00:00<00:00, 62.37it/s]

Finished epoch 5 of 50: Train Loss = 1.962    Train Acc = 0.254    Test Acc = 0.272

100%|██████████| 78/78 [00:03<00:00, 24.38it/s]

100%|██████████| 16/16 [00:00<00:00, 61.30it/s]

Finished epoch 6 of 50: Train Loss = 1.900    Train Acc = 0.275    Test Acc = 0.287

100%|██████████| 78/78 [00:03<00:00, 25.62it/s]

100%|██████████| 16/16 [00:00<00:00, 62.28it/s]

Finished epoch 7 of 50: Train Loss = 1.857    Train Acc = 0.293    Test Acc = 0.317

100%|██████████| 78/78 [00:03<00:00, 24.48it/s]

100%|██████████| 16/16 [00:00<00:00, 62.28it/s]

Finished epoch 8 of 50: Train Loss = 1.819    Train Acc = 0.309    Test Acc = 0.315

100%|██████████| 78/78 [00:03<00:00, 24.64it/s]

100%|██████████| 16/16 [00:00<00:00, 63.05it/s]

Finished epoch 9 of 50: Train Loss = 1.778    Train Acc = 0.330    Test Acc = 0.349

100%|██████████| 78/78 [00:03<00:00, 24.80it/s]

100%|██████████| 16/16 [00:00<00:00, 63.01it/s]

Finished epoch 10 of 50: Train Loss = 1.745    Train Acc = 0.345    Test Acc = 0.364

100%|██████████| 78/78 [00:03<00:00, 25.00it/s]

100%|██████████| 16/16 [00:00<00:00, 60.82it/s]

Finished epoch 11 of 50: Train Loss = 1.710    Train Acc = 0.361    Test Acc = 0.371

100%|██████████| 78/78 [00:03<00:00, 24.87it/s]

100%|██████████| 16/16 [00:00<00:00, 53.65it/s]

Finished epoch 12 of 50: Train Loss = 1.684    Train Acc = 0.378    Test Acc = 0.396

100%|██████████| 78/78 [00:03<00:00, 25.35it/s]

100%|██████████| 16/16 [00:00<00:00, 62.56it/s]

Finished epoch 13 of 50: Train Loss = 1.656    Train Acc = 0.380    Test Acc = 0.392

100%|██████████| 78/78 [00:03<00:00, 24.80it/s]

100%|██████████| 16/16 [00:00<00:00, 63.11it/s]

Finished epoch 14 of 50: Train Loss = 1.641    Train Acc = 0.396    Test Acc = 0.388

100%|██████████| 78/78 [00:03<00:00, 24.26it/s]

100%|██████████| 16/16 [00:00<00:00, 62.16it/s]

Finished epoch 15 of 50: Train Loss = 1.619    Train Acc = 0.396    Test Acc = 0.385

100%|██████████| 78/78 [00:03<00:00, 25.28it/s]

100%|██████████| 16/16 [00:00<00:00, 62.73it/s]

Finished epoch 16 of 50: Train Loss = 1.612    Train Acc = 0.401    Test Acc = 0.407

100%|██████████| 78/78 [00:03<00:00, 24.34it/s]

100%|██████████| 16/16 [00:00<00:00, 63.64it/s]

Finished epoch 17 of 50: Train Loss = 1.598    Train Acc = 0.415    Test Acc = 0.368

100%|██████████| 78/78 [00:03<00:00, 24.68it/s]

100%|██████████| 16/16 [00:00<00:00, 63.06it/s]

Finished epoch 18 of 50: Train Loss = 1.580    Train Acc = 0.421    Test Acc = 0.399

100%|██████████| 78/78 [00:03<00:00, 25.03it/s]

100%|██████████| 16/16 [00:00<00:00, 63.96it/s]

Finished epoch 19 of 50: Train Loss = 1.580    Train Acc = 0.417    Test Acc = 0.417

100%|██████████| 78/78 [00:03<00:00, 24.86it/s]

100%|██████████| 16/16 [00:00<00:00, 62.02it/s]

Finished epoch 20 of 50: Train Loss = 1.560    Train Acc = 0.421    Test Acc = 0.420

100%|██████████| 78/78 [00:03<00:00, 24.74it/s]

100%|██████████| 16/16 [00:00<00:00, 52.94it/s]

Finished epoch 21 of 50: Train Loss = 1.533 Train Acc = 0.438 Test Acc = 0.423

100%|██████████| 78/78 [00:03<00:00, 25.20it/s]

100%|██████████| 16/16 [00:00<00:00, 62.75it/s]

Finished epoch 22 of 50: Train Loss = 1.513 Train Acc = 0.447 Test Acc = 0.427

100%|██████████| 78/78 [00:03<00:00, 24.83it/s]

100%|██████████| 16/16 [00:00<00:00, 61.47it/s]

Finished epoch 23 of 50: Train Loss = 1.496 Train Acc = 0.452 Test Acc = 0.433

100%|██████████| 78/78 [00:03<00:00, 24.43it/s]

100%|██████████| 16/16 [00:00<00:00, 61.96it/s]

Finished epoch 24 of 50: Train Loss = 1.494 Train Acc = 0.454 Test Acc = 0.433

100%|██████████| 78/78 [00:03<00:00, 25.22it/s]

100%|██████████| 16/16 [00:00<00:00, 61.62it/s]

Finished epoch 25 of 50: Train Loss = 1.470 Train Acc = 0.461 Test Acc = 0.419

100%|██████████| 78/78 [00:03<00:00, 24.52it/s]

100%|██████████| 16/16 [00:00<00:00, 54.82it/s]

Finished epoch 26 of 50: Train Loss = 1.451 Train Acc = 0.473 Test Acc = 0.436

100%|██████████| 78/78 [00:03<00:00, 24.80it/s]

100%|██████████| 16/16 [00:00<00:00, 61.61it/s]

Finished epoch 27 of 50: Train Loss = 1.433 Train Acc = 0.475 Test Acc = 0.443

100%|██████████| 78/78 [00:03<00:00, 25.08it/s]

100%|██████████| 16/16 [00:00<00:00, 58.62it/s]

Finished epoch 28 of 50: Train Loss = 1.419 Train Acc = 0.489 Test Acc = 0.446

100%|██████████| 78/78 [00:03<00:00, 24.89it/s]

100%|██████████| 16/16 [00:00<00:00, 55.37it/s]

Finished epoch 29 of 50: Train Loss = 1.416 Train Acc = 0.486 Test Acc = 0.446

100%|██████████| 78/78 [00:03<00:00, 24.44it/s]

100%|██████████| 16/16 [00:00<00:00, 61.08it/s]

Finished epoch 30 of 50: Train Loss = 1.413 Train Acc = 0.492 Test Acc = 0.467

100%|██████████| 78/78 [00:03<00:00, 25.03it/s]

100%|██████████| 16/16 [00:00<00:00, 57.89it/s]

Finished epoch 31 of 50: Train Loss = 1.394 Train Acc = 0.494 Test Acc = 0.458

100%|██████████| 78/78 [00:03<00:00, 25.04it/s]

100%|██████████| 16/16 [00:00<00:00, 57.32it/s]

Finished epoch 32 of 50: Train Loss = 1.396 Train Acc = 0.502 Test Acc = 0.450

100%|██████████| 78/78 [00:03<00:00, 23.51it/s]

100%|██████████| 16/16 [00:00<00:00, 54.95it/s]



Finished epoch 33 of 50: Train Loss = 1.367    Train Acc = 0.504    Test Acc = 0.465

100%|██████████| 78/78 [00:03<00:00, 23.32it/s]

100%|██████████| 16/16 [00:00<00:00, 51.17it/s]

Finished epoch 34 of 50: Train Loss = 1.354    Train Acc = 0.509    Test Acc = 0.451

100%|██████████| 78/78 [00:03<00:00, 23.95it/s]

100%|██████████| 16/16 [00:00<00:00, 47.46it/s]

Finished epoch 35 of 50: Train Loss = 1.343    Train Acc = 0.516    Test Acc = 0.456

100%|██████████| 78/78 [00:03<00:00, 23.72it/s]

100%|██████████| 16/16 [00:00<00:00, 59.20it/s]

Finished epoch 36 of 50: Train Loss = 1.336    Train Acc = 0.519    Test Acc = 0.454

100%|██████████| 78/78 [00:03<00:00, 23.96it/s]

100%|██████████| 16/16 [00:00<00:00, 62.19it/s]

Finished epoch 37 of 50: Train Loss = 1.344    Train Acc = 0.514    Test Acc = 0.463

100%|██████████| 78/78 [00:03<00:00, 23.58it/s]

100%|██████████| 16/16 [00:00<00:00, 54.76it/s]

Finished epoch 38 of 50: Train Loss = 1.325    Train Acc = 0.521    Test Acc = 0.470

100%|██████████| 78/78 [00:03<00:00, 24.31it/s]

100%|██████████| 16/16 [00:00<00:00, 65.93it/s]

Finished epoch 39 of 50: Train Loss = 1.314    Train Acc = 0.529    Test Acc = 0.479

100%|██████████| 78/78 [00:03<00:00, 23.53it/s]

100%|██████████| 16/16 [00:00<00:00, 57.81it/s]

Finished epoch 40 of 50: Train Loss = 1.304    Train Acc = 0.526    Test Acc = 0.471

100%|██████████| 78/78 [00:03<00:00, 23.23it/s]

100%|██████████| 16/16 [00:00<00:00, 60.06it/s]

Finished epoch 41 of 50: Train Loss = 1.303    Train Acc = 0.527    Test Acc = 0.483

100%|██████████| 78/78 [00:03<00:00, 23.19it/s]

100%|██████████| 16/16 [00:00<00:00, 59.77it/s]

Finished epoch 42 of 50: Train Loss = 1.279    Train Acc = 0.541    Test Acc = 0.479

100%|██████████| 78/78 [00:03<00:00, 23.88it/s]

100%|██████████| 16/16 [00:00<00:00, 59.78it/s]

Finished epoch 43 of 50: Train Loss = 1.273    Train Acc = 0.537    Test Acc = 0.479

100%|██████████| 78/78 [00:03<00:00, 23.89it/s]

100%|██████████| 16/16 [00:00<00:00, 56.93it/s]

Finished epoch 44 of 50: Train Loss = 1.274    Train Acc = 0.543    Test Acc = 0.474

100%|██████████| 78/78 [00:03<00:00, 25.69it/s]

100%|██████████| 16/16 [00:00<00:00, 54.08it/s]

Finished epoch 45 of 50: Train Loss = 1.259    Train Acc = 0.548    Test Acc = 0.488

100%|██████████| 78/78 [00:03<00:00, 24.04it/s]

100%|██████████| 16/16 [00:00<00:00, 58.26it/s]

Finished epoch 46 of 50: Train Loss = 1.244    Train Acc = 0.553    Test Acc = 0.497

100%|██████████| 78/78 [00:03<00:00, 23.71it/s]

100%|██████████| 16/16 [00:00<00:00, 59.21it/s]

Finished epoch 47 of 50: Train Loss = 1.248    Train Acc = 0.553    Test Acc = 0.478

100%|██████████| 78/78 [00:03<00:00, 25.53it/s]

100%|██████████| 16/16 [00:00<00:00, 65.82it/s]

Finished epoch 48 of 50: Train Loss = 1.234    Train Acc = 0.560    Test Acc = 0.493

100%|██████████| 78/78 [00:02<00:00, 26.59it/s]

100%|██████████| 16/16 [00:00<00:00, 66.46it/s]

Finished epoch 49 of 50: Train Loss = 1.226    Train Acc = 0.556    Test Acc = 0.498


## Evaluate the model

```
In [27]: plt.plot(hist_tr_acc, label='train accuracy')
plt.plot(hist_te_acc, label='test_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.3, 0.7])
plt.legend(loc='lower right')

acc = eval(model, loader_test)
print('Test accuracy:', acc)
```

100%|██████████| 16/16 [00:00<00:00, 61.76it/s]

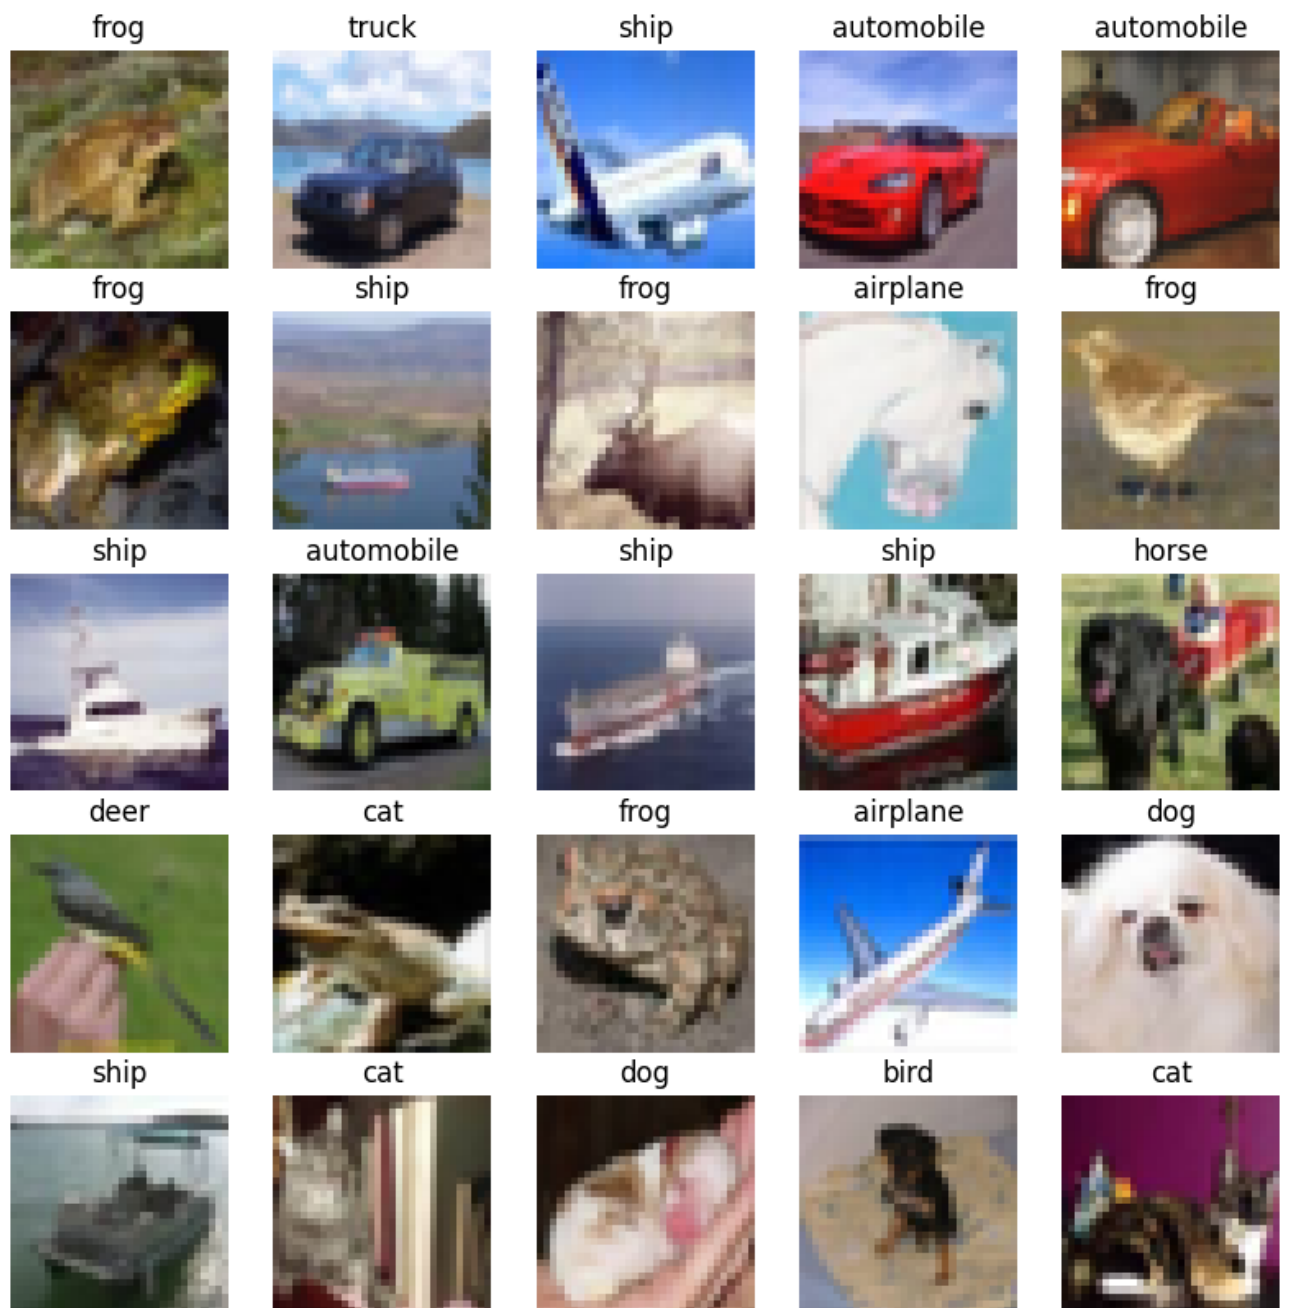
Test accuracy: 0.49843749962747097



Now lets display a bunch of predictions of the model


```
In [28]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                        'dog', 'frog', 'horse', 'ship', 'truck']

plt.figure(figsize=(10,10))
for i in range(25):
    img, lbl = db_train[i]
    with torch.no_grad():
        pred = model(img[None])[0]
        pred = pred.argmax()
    plt.subplot(5,5,i+1)
    plt.imshow(img.permute(1, 2, 0))
    plt.title(class_names[pred])
    plt.axis(False)
plt.show()
```



Now let's see what the model learned at the first layer. Do you recognize any of these filters?

```
In [29]: W = (model[0].weight - model[0].weight.min()) / (model[0].weight.max() - model[0].weight.min())
for i in range(32):
    plt.subplot(8, 4, i+1)
    plt.imshow(W[i].permute(1, 2, 0).detach())
    plt.axis(False)
```



In [ ]:

```
In [4]: import matplotlib.pyplot as plt
import numpy as np

import torch
import torch.nn as nn
from torchvision import datasets
from torchvision import transforms

import ssl
ssl._create_default_https_context = ssl._create_unverified_context

transform = transforms.Compose([
    transforms.ToTensor(),
])
db_train = datasets.CIFAR10(root='./', train=True, transform=transform, download=True)
db_test = datasets.CIFAR10(root='./', train=False, transform=transform, download=True)
def subsample(db, n):
    idx = np.random.permutation(len(db))[:n]
    db.data = db.data[idx]
    db.targets = np.array(db.targets)[idx]
    return db

db_train = subsample(db_train, 5000)
db_test = subsample(db_test, 1000)

img, y = db_train[0]
```

```
print(img.shape, img.min(), img.max(), y)
```

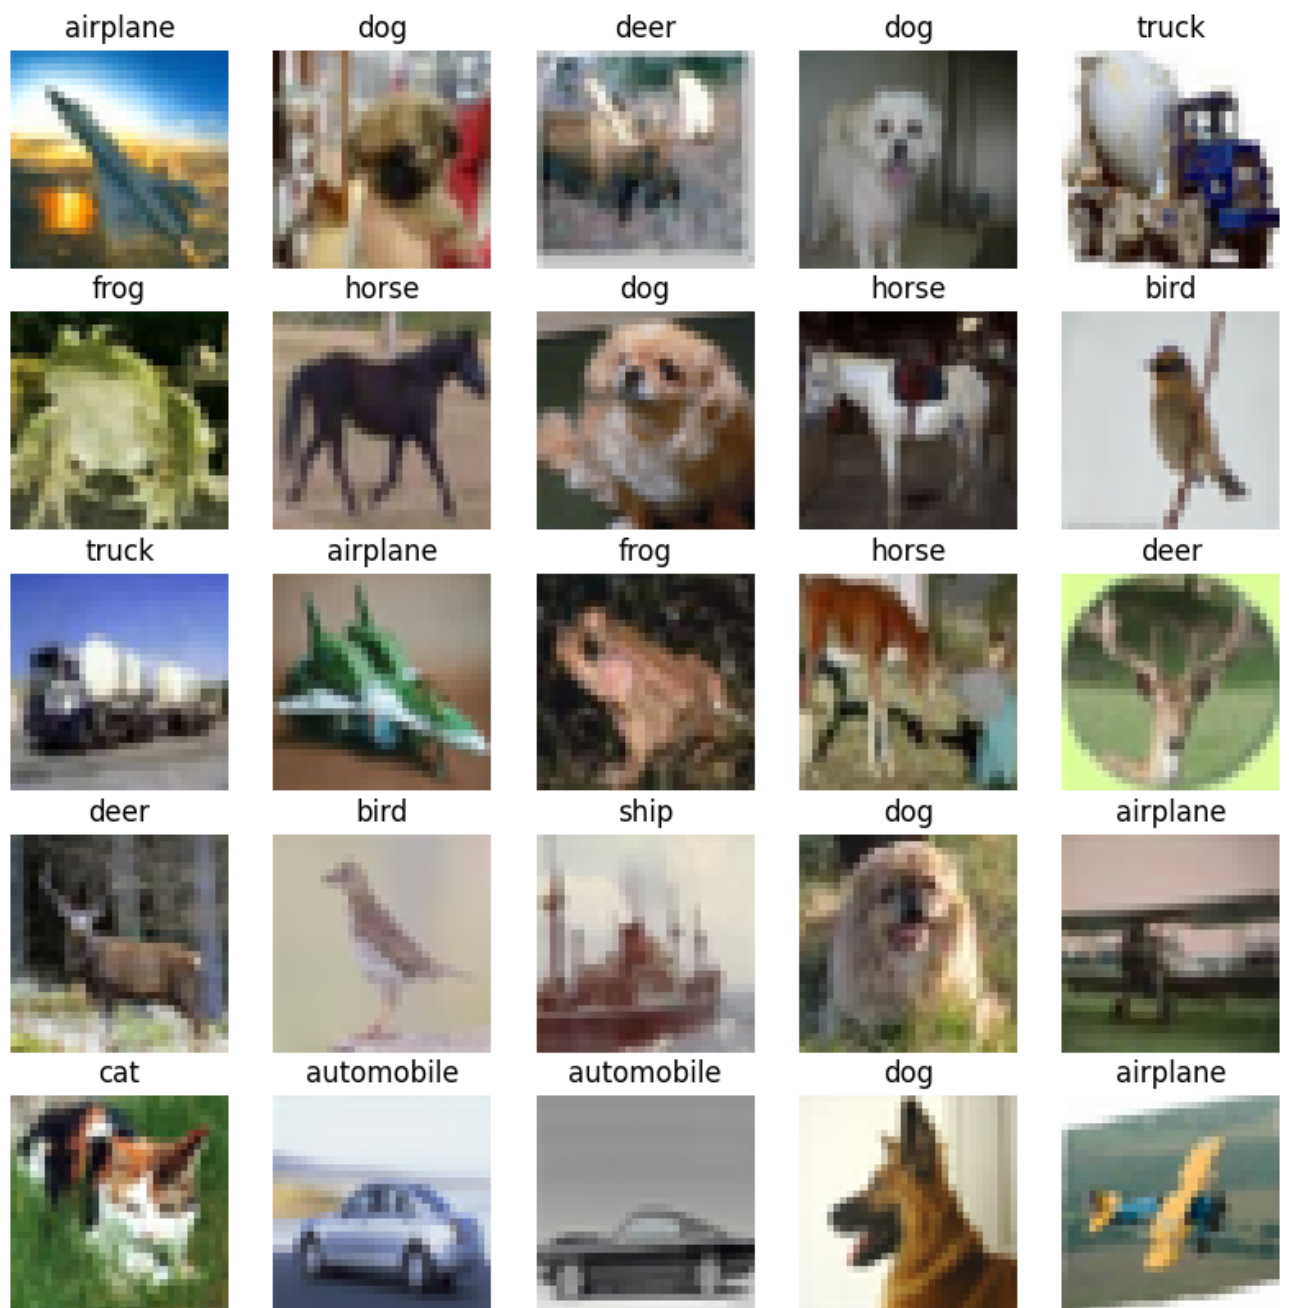
Files already downloaded and verified

Files already downloaded and verified

torch.Size([3, 32, 32]) tensor(0.0078) tensor(1.) 0

```
In [5]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
                        'dog', 'frog', 'horse', 'ship', 'truck']
```

```
plt.figure(figsize=(10,10))  
for i in range(25):  
    img, lbl = db_train[i]  
    plt.subplot(5,5,i+1)  
    plt.imshow(img.permute(1, 2, 0))  
    plt.title(class_names[lbl])  
    plt.axis(False)  
plt.show()
```



```
In [19]: model = nn.Sequential(
    nn.Conv2d(3, 32, kernel_size=(5, 5)),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=(2, 2), stride=(2, 2)),
    nn.Conv2d(32, 64, kernel_size=(3, 3)),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=(2, 2), stride=(2, 2)),
    nn.Conv2d(64, 64, kernel_size=(3, 3)),
    nn.ReLU(),
)
dummy_img = torch.randn(1, 3, 32, 32)
print(model(dummy_img).shape)
```

```
torch.Size([1, 64, 4, 4])
```

```
In [8]: from torch.utils.data import DataLoader
```

```

import tqdm
device = 'cuda' if torch.cuda.is_available() else 'cpu'
def init_weights(model):
    for m in model.modules():
        if isinstance(m, (nn.Conv2d, nn.Linear)):
            nn.init.normal_(m.weight, mean=0, std=0.005)

def accuracy(pred, labels):
    return (pred.argmax(-1) == labels).float().mean()

@torch.no_grad()
def eval(model, loader):
    accuracies = []
    for images, labels in tqdm.tqdm(loader):
        # Prepare data
        images = images.to(device)
        labels = labels.to(device)

        # Compute predictions, loss and accuracy
        pred = model(images)
        acc = accuracy(pred, labels)
        accuracies.append(acc.detach().item())
    return np.mean(accuracies)

def fit_one_epoch(model, opt, loader):
    losses, accuracies = [], []
    for images, labels in tqdm.tqdm(loader):
        # Prepare data
        images = images.to(device)
        labels = labels.to(device)

        # Compute predictions, loss and accuracy
        pred = model(images)
        l = loss(pred, labels)
        acc = accuracy(pred, labels)

        # Compute gradients and update the model
        opt.zero_grad()
        l.backward()
        opt.step()

        losses.append(l.detach().item())
        accuracies.append(acc.detach().item())
    return np.mean(losses), np.mean(accuracies)

def fit(model, loader_train, loader_test, epochs=10):
    opt = torch.optim.Adam(model.parameters(), lr=0.0005)
    hist_tr_loss, hist_tr_acc, hist_te_acc = [], [], []
    for epoch in range(epochs):
        tr_l, tr_acc = fit_one_epoch(model, opt, loader_train)
        te_acc = eval(model, loader_test)

```



```

        print(f"Finished epoch {epoch} of {epochs}: Train Loss = {tr_l:.3f}")
        hist_tr_loss.append(tr_l)
        hist_tr_acc.append(tr_acc)
        hist_te_acc.append(te_acc)
    return hist_tr_loss, hist_tr_acc, hist_te_acc

loader_train = DataLoader(db_train, batch_size=64, shuffle=True, drop_last=True)
loader_test = DataLoader(db_test, batch_size=64, shuffle=False)
loss = nn.CrossEntropyLoss()
model = model.to(device)
init_weights(model)
hist_tr_loss, hist_tr_acc, hist_te_acc = fit(model, loader_train, loader_test)

```

100%|██████████| 78/78 [00:03<00:00, 25.35it/s]

100%|██████████| 16/16 [00:00<00:00, 66.58it/s]

Finished epoch 0 of 30: Train Loss = 2.276    Train Acc = 0.122    Test Acc = 0.194

100%|██████████| 78/78 [00:02<00:00, 26.84it/s]

100%|██████████| 16/16 [00:00<00:00, 67.87it/s]

Finished epoch 1 of 30: Train Loss = 2.106    Train Acc = 0.207    Test Acc = 0.259

100%|██████████| 78/78 [00:02<00:00, 26.96it/s]

100%|██████████| 16/16 [00:00<00:00, 68.71it/s]

Finished epoch 2 of 30: Train Loss = 2.040    Train Acc = 0.242    Test Acc = 0.271

100%|██████████| 78/78 [00:03<00:00, 25.85it/s]

100%|██████████| 16/16 [00:00<00:00, 65.47it/s]

Finished epoch 3 of 30: Train Loss = 2.017    Train Acc = 0.252    Test Acc = 0.250

100%|██████████| 78/78 [00:02<00:00, 26.44it/s]

100%|██████████| 16/16 [00:00<00:00, 67.15it/s]

Finished epoch 4 of 30: Train Loss = 1.998    Train Acc = 0.258    Test Acc = 0.288

100%|██████████| 78/78 [00:02<00:00, 26.08it/s]

100%|██████████| 16/16 [00:00<00:00, 67.55it/s]

Finished epoch 5 of 30: Train Loss = 1.964    Train Acc = 0.267    Test Acc = 0.273

100%|██████████| 78/78 [00:02<00:00, 26.27it/s]

100%|██████████| 16/16 [00:00<00:00, 63.48it/s]

Finished epoch 6 of 30: Train Loss = 1.921    Train Acc = 0.289    Test Acc = 0.291

100%|██████████| 78/78 [00:03<00:00, 25.92it/s]

100%|██████████| 16/16 [00:00<00:00, 66.49it/s]

Finished epoch 7 of 30: Train Loss = 1.871    Train Acc = 0.310    Test Acc = 0.321

100%|██████████| 78/78 [00:02<00:00, 26.19it/s]

100%|██████████| 16/16 [00:00<00:00, 66.55it/s]

Finished epoch 8 of 30: Train Loss = 1.830    Train Acc = 0.332    Test Acc = 0.334

100%|██████████| 78/78 [00:02<00:00, 26.13it/s]

100%|██████████| 16/16 [00:00<00:00, 66.28it/s]

Finished epoch 9 of 30: Train Loss = 1.779 Train Acc = 0.350 Test Acc = 0.352

100%|██████████| 78/78 [00:02<00:00, 26.23it/s]

100%|██████████| 16/16 [00:00<00:00, 66.51it/s]

Finished epoch 10 of 30: Train Loss = 1.736 Train Acc = 0.363 Test Acc = 0.346

100%|██████████| 78/78 [00:02<00:00, 26.16it/s]

100%|██████████| 16/16 [00:00<00:00, 62.23it/s]

Finished epoch 11 of 30: Train Loss = 1.711 Train Acc = 0.374 Test Acc = 0.370

100%|██████████| 78/78 [00:02<00:00, 26.62it/s]

100%|██████████| 16/16 [00:00<00:00, 65.97it/s]

Finished epoch 12 of 30: Train Loss = 1.679 Train Acc = 0.385 Test Acc = 0.392

100%|██████████| 78/78 [00:03<00:00, 25.88it/s]

100%|██████████| 16/16 [00:00<00:00, 66.01it/s]

Finished epoch 13 of 30: Train Loss = 1.675 Train Acc = 0.393 Test Acc = 0.383

100%|██████████| 78/78 [00:03<00:00, 25.80it/s]

100%|██████████| 16/16 [00:00<00:00, 66.68it/s]

Finished epoch 14 of 30: Train Loss = 1.661 Train Acc = 0.402 Test Acc = 0.370

100%|██████████| 78/78 [00:03<00:00, 24.99it/s]

100%|██████████| 16/16 [00:00<00:00, 64.11it/s]

Finished epoch 15 of 30: Train Loss = 1.629 Train Acc = 0.413 Test Acc = 0.393

100%|██████████| 78/78 [00:03<00:00, 25.80it/s]

100%|██████████| 16/16 [00:00<00:00, 66.26it/s]

Finished epoch 16 of 30: Train Loss = 1.611 Train Acc = 0.420 Test Acc = 0.396

100%|██████████| 78/78 [00:02<00:00, 26.33it/s]

100%|██████████| 16/16 [00:00<00:00, 67.87it/s]

Finished epoch 17 of 30: Train Loss = 1.588 Train Acc = 0.422 Test Acc = 0.413

100%|██████████| 78/78 [00:02<00:00, 26.30it/s]

100%|██████████| 16/16 [00:00<00:00, 64.85it/s]

Finished epoch 18 of 30: Train Loss = 1.562 Train Acc = 0.437 Test Acc = 0.411

100%|██████████| 78/78 [00:02<00:00, 26.04it/s]

100%|██████████| 16/16 [00:00<00:00, 67.17it/s]

Finished epoch 19 of 30: Train Loss = 1.551 Train Acc = 0.438 Test Acc = 0.413

100%|██████████| 78/78 [00:02<00:00, 26.18it/s]

100%|██████████| 16/16 [00:00<00:00, 57.36it/s]

Finished epoch 20 of 30: Train Loss = 1.552 Train Acc = 0.446 Test Acc = 0.402

100%|██████████| 78/78 [00:02<00:00, 26.38it/s]

100%|██████████| 16/16 [00:00<00:00, 65.66it/s]

```

Finished epoch 21 of 30: Train Loss = 1.512   Train Acc = 0.453   Test Acc = 0.431
100%|██████████| 78/78 [00:03<00:00, 25.81it/s]
100%|██████████| 16/16 [00:00<00:00, 67.08it/s]
Finished epoch 22 of 30: Train Loss = 1.509   Train Acc = 0.455   Test Acc = 0.428
100%|██████████| 78/78 [00:03<00:00, 25.81it/s]
100%|██████████| 16/16 [00:00<00:00, 66.63it/s]
Finished epoch 23 of 30: Train Loss = 1.485   Train Acc = 0.461   Test Acc = 0.431
100%|██████████| 78/78 [00:02<00:00, 26.09it/s]
100%|██████████| 16/16 [00:00<00:00, 58.76it/s]
Finished epoch 24 of 30: Train Loss = 1.484   Train Acc = 0.468   Test Acc = 0.423
100%|██████████| 78/78 [00:03<00:00, 25.56it/s]
100%|██████████| 16/16 [00:00<00:00, 65.11it/s]
Finished epoch 25 of 30: Train Loss = 1.468   Train Acc = 0.468   Test Acc = 0.440
100%|██████████| 78/78 [00:02<00:00, 26.02it/s]
100%|██████████| 16/16 [00:00<00:00, 64.42it/s]
Finished epoch 26 of 30: Train Loss = 1.462   Train Acc = 0.472   Test Acc = 0.430
100%|██████████| 78/78 [00:03<00:00, 25.89it/s]
100%|██████████| 16/16 [00:00<00:00, 64.70it/s]
Finished epoch 27 of 30: Train Loss = 1.429   Train Acc = 0.486   Test Acc = 0.436
100%|██████████| 78/78 [00:02<00:00, 26.07it/s]
100%|██████████| 16/16 [00:00<00:00, 60.38it/s]
Finished epoch 28 of 30: Train Loss = 1.409   Train Acc = 0.493   Test Acc = 0.434
100%|██████████| 78/78 [00:03<00:00, 26.00it/s]
100%|██████████| 16/16 [00:00<00:00, 56.90it/s]
Finished epoch 29 of 30: Train Loss = 1.413   Train Acc = 0.493   Test Acc = 0.453

```

```

In [9]: plt.figure(figsize=(15,8))
plt.subplot(1, 2, 1)
plt.plot(hist_tr_acc, label='train accuracy')
plt.plot(hist_te_acc, label='test accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

plt.subplot(1, 2, 2)
plt.plot(hist_tr_loss, label='train loss')
plt.xlabel('Epoch')
plt.ylabel('Training Loss')
plt.legend(loc='lower right')
plt.tight_layout()


acc = eval(model, loader_test)

```

```
print(f'Test accuracy: {acc*100:.2f}%')
```

100%|██████████| 16/16 [00:00<00:00, 54.15it/s]

Test accuracy: 45.29%



```
In [10]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
                        'dog', 'frog', 'horse', 'ship', 'truck']  
model = model.to(device)  
  
plt.figure(figsize=(20,20))  
for imgs, lbl in loader_test:  
    with torch.no_grad():  
        imgs = imgs.to(device)  
        preds = model(imgs)  
        preds = preds.argmax(1)  
        for i in range(25):  
            plt.subplot(5,5,i+1)  
            plt.imshow(imgs[i].cpu().permute(1, 2, 0))  
            plt.title(class_names[preds[i]])  
            # plt.axis(False)  
plt.show()
```



```
In [15]: class MyBatchNorm(nn.Module):
def __init__(self, dims):
    super().__init__()
    self.alpha = nn.Parameter(torch.ones(1, dims, 1, 1))
    self.beta = nn.Parameter(torch.zeros(1, dims, 1, 1))

def forward(self, x):
    # Compute mean and standard deviation across batch, height, and width
    mean = x.mean(dim=(0, 2, 3), keepdim=True)
    std = x.std(dim=(0, 2, 3), keepdim=True, unbiased=False) # Unbiased

    # Normalize input
    x_norm = (x - mean) / (std)

    # Scale and shift
    x_out = self.alpha * x_norm + self.beta
```

```

        return x_out

model = nn.Sequential(
    nn.Conv2d(3, 32, kernel_size=(5, 5)),
    MyBatchNorm(32),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=(2, 2), stride=(2, 2)),
    nn.Conv2d(32, 64, kernel_size=(3, 3)),
    MyBatchNorm(64),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=(2, 2), stride=(2, 2)),
    nn.Conv2d(64, 64, kernel_size=(3, 3)),
    MyBatchNorm(64),
    nn.ReLU(),
    nn.Flatten(),
    nn.Linear(1024, 10)
)

# Check the model
dummy_image = torch.randn(1, 3, 32, 32)
dummy_output = model(dummy_image)
print(f"{str(dummy_image.shape)} -> {str(dummy_output.shape)}")

torch.Size([1, 3, 32, 32]) -> torch.Size([1, 10])

```

```

In [ ]: # Train the model
init_weights(model)
model = model.to(device)
hist_tr_loss, hist_tr_acc, hist_te_acc = fit(model, loader_train, loader_test)

# Plot training history
plt.subplot(1, 2, 1)
plt.plot(hist_tr_acc, label='train accuracy')
plt.plot(hist_te_acc, label='test accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
# plt.ylim([0.3, 1])
plt.legend(loc='lower right')

plt.subplot(1, 2, 2)
plt.plot(hist_tr_loss, label='train loss')
plt.xlabel('Epoch')
plt.ylabel('Training Loss')
plt.legend(loc='lower right')
plt.tight_layout()

# Evaluate
acc = eval(model, loader_test)
print(f'Test accuracy: {acc*100:.2f}%')

```

```

In [16]: def display_transform(transform, image_idx=10):
        # The CIFAR dataset uses db.transform to transform each image.

```

```

# We can change the data augmentation by assigning our new transformatio
# See __getitem__ in https://pytorch.org/vision/stable/_modules/torchvis
db_train.transform = transform
plt.figure(figsize=(20,6))
for i in range(5):
    img, lbl = db_train[image_idx]
    plt.subplot(1,5,1+i)
    plt.imshow(img.permute(1, 2, 0))
    plt.axis(False)
plt.show()

print("Transform 1")
transform1 = transforms.Compose([
    transforms.RandomCrop(size=[32,32], padding=3),
    transforms.ToTensor()
])
display_transform(transform1, image_idx=0)


print("Transform 2")
transform2 = transforms.Compose([
    transforms.RandomHorizontalFlip(0.5),
    transforms.ToTensor()
])
display_transform(transform2, image_idx=0)

print("Transform 3")
transform3 = transforms.Compose([
    transforms.ColorJitter(brightness=0.5),
    transforms.ToTensor()
])
display_transform(transform3, image_idx=0)





print("Transform 4")
transform4 = transforms.Compose([
    transforms.ColorJitter(hue=0.5),
    transforms.ToTensor()
])
display_transform(transform4, image_idx=0)

```


Transform 1




Transform 2



Transform 3



Transform 4



```
In [18]: # Define new transform
db_train.transform = transforms.Compose([
    transforms.RandomRotation(40),
    transforms.ToTensor()
])

loader_train = DataLoader(db_train, batch_size=64, shuffle=True, drop_last=True)
loader_test = DataLoader(db_test, batch_size=64, shuffle=False)
loss = nn.CrossEntropyLoss()
init_weights(model)
hist_tr_loss, hist_tr_acc, hist_te_acc = fit(model, loader_train, loader_test)

# Plot training history
plt.subplot(1, 2, 1)
plt.plot(hist_tr_acc, label='train accuracy')
plt.plot(hist_te_acc, label='test accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
# plt.ylim([0.3, 1])
plt.legend(loc='lower right')

plt.subplot(1, 2, 2)
plt.plot(hist_tr_loss, label='train loss')
plt.xlabel('Epoch')
plt.ylabel('Training Loss')
plt.legend(loc='lower right')
plt.tight_layout()

# Evaluate
acc = eval(model, loader_test)
```



```
print(f'Test accuracy: {acc*100:.2f}%')
```

100%|██████████| 78/78 [00:05<00:00, 14.75it/s]

100%|██████████| 16/16 [00:00<00:00, 33.28it/s]

Finished epoch 0 of 30: Train Loss = 1.839 Train Acc = 0.324 Test Acc = 0.373

100%|██████████| 78/78 [00:05<00:00, 14.96it/s]

100%|██████████| 16/16 [00:00<00:00, 35.98it/s]

Finished epoch 1 of 30: Train Loss = 1.596 Train Acc = 0.413 Test Acc = 0.443

100%|██████████| 78/78 [00:05<00:00, 14.80it/s]

100%|██████████| 16/16 [00:00<00:00, 38.02it/s]

Finished epoch 2 of 30: Train Loss = 1.482 Train Acc = 0.463 Test Acc = 0.466

100%|██████████| 78/78 [00:05<00:00, 15.01it/s]

100%|██████████| 16/16 [00:00<00:00, 33.79it/s]

Finished epoch 3 of 30: Train Loss = 1.416 Train Acc = 0.488 Test Acc = 0.529

100%|██████████| 78/78 [00:05<00:00, 14.74it/s]

100%|██████████| 16/16 [00:00<00:00, 36.33it/s]

Finished epoch 4 of 30: Train Loss = 1.333 Train Acc = 0.517 Test Acc = 0.496

100%|██████████| 78/78 [00:05<00:00, 14.69it/s]

100%|██████████| 16/16 [00:00<00:00, 36.04it/s]

Finished epoch 5 of 30: Train Loss = 1.297 Train Acc = 0.529 Test Acc = 0.516

100%|██████████| 78/78 [00:05<00:00, 14.98it/s]

100%|██████████| 16/16 [00:00<00:00, 35.64it/s]

Finished epoch 6 of 30: Train Loss = 1.237 Train Acc = 0.557 Test Acc = 0.580

100%|██████████| 78/78 [00:05<00:00, 14.59it/s]

100%|██████████| 16/16 [00:00<00:00, 36.86it/s]

Finished epoch 7 of 30: Train Loss = 1.206 Train Acc = 0.578 Test Acc = 0.557

100%|██████████| 78/78 [00:05<00:00, 14.61it/s]

100%|██████████| 16/16 [00:00<00:00, 36.19it/s]

Finished epoch 8 of 30: Train Loss = 1.162 Train Acc = 0.588 Test Acc = 0.546

100%|██████████| 78/78 [00:05<00:00, 14.76it/s]

100%|██████████| 16/16 [00:00<00:00, 36.92it/s]

Finished epoch 9 of 30: Train Loss = 1.129 Train Acc = 0.606 Test Acc = 0.572

100%|██████████| 78/78 [00:05<00:00, 14.65it/s]

100%|██████████| 16/16 [00:00<00:00, 36.49it/s]

Finished epoch 10 of 30: Train Loss = 1.116 Train Acc = 0.600 Test Acc = 0.587

100%|██████████| 78/78 [00:05<00:00, 14.75it/s]

100%|██████████| 16/16 [00:00<00:00, 31.67it/s]

Finished epoch 11 of 30: Train Loss = 1.073 Train Acc = 0.624 Test Acc = 0.570

100%|██████████| 78/78 [00:05<00:00, 14.91it/s]  
100%|██████████| 16/16 [00:00<00:00, 36.19it/s]  
Finished epoch 12 of 30: Train Loss = 1.027    Train Acc = 0.642    Test Acc = 0.593

100%|██████████| 78/78 [00:05<00:00, 14.73it/s]  
100%|██████████| 16/16 [00:00<00:00, 37.50it/s]  
Finished epoch 13 of 30: Train Loss = 0.998    Train Acc = 0.647    Test Acc = 0.599

100%|██████████| 78/78 [00:05<00:00, 14.59it/s]  
100%|██████████| 16/16 [00:00<00:00, 35.51it/s]  
Finished epoch 14 of 30: Train Loss = 0.986    Train Acc = 0.655    Test Acc = 0.581

100%|██████████| 78/78 [00:05<00:00, 14.60it/s]  
100%|██████████| 16/16 [00:00<00:00, 35.17it/s]  
Finished epoch 15 of 30: Train Loss = 0.961    Train Acc = 0.668    Test Acc = 0.602

100%|██████████| 78/78 [00:05<00:00, 14.58it/s]  
100%|██████████| 16/16 [00:00<00:00, 35.83it/s]  
Finished epoch 16 of 30: Train Loss = 0.919    Train Acc = 0.678    Test Acc = 0.579

100%|██████████| 78/78 [00:05<00:00, 14.71it/s]  
100%|██████████| 16/16 [00:00<00:00, 34.45it/s]  
Finished epoch 17 of 30: Train Loss = 0.917    Train Acc = 0.687    Test Acc = 0.586

100%|██████████| 78/78 [00:05<00:00, 14.53it/s]  
100%|██████████| 16/16 [00:00<00:00, 26.23it/s]  
Finished epoch 18 of 30: Train Loss = 0.891    Train Acc = 0.690    Test Acc = 0.607

100%|██████████| 78/78 [00:05<00:00, 14.65it/s]  
100%|██████████| 16/16 [00:00<00:00, 35.51it/s]  
Finished epoch 19 of 30: Train Loss = 0.885    Train Acc = 0.693    Test Acc = 0.615

100%|██████████| 78/78 [00:05<00:00, 14.77it/s]  
100%|██████████| 16/16 [00:00<00:00, 36.41it/s]  
Finished epoch 20 of 30: Train Loss = 0.834    Train Acc = 0.716    Test Acc = 0.610

100%|██████████| 78/78 [00:05<00:00, 14.60it/s]  
100%|██████████| 16/16 [00:00<00:00, 35.90it/s]  
Finished epoch 21 of 30: Train Loss = 0.829    Train Acc = 0.711    Test Acc = 0.576

100%|██████████| 78/78 [00:05<00:00, 14.62it/s]  
100%|██████████| 16/16 [00:00<00:00, 36.07it/s]  
Finished epoch 22 of 30: Train Loss = 0.821    Train Acc = 0.714    Test Acc = 0.603

100%|██████████| 78/78 [00:05<00:00, 14.69it/s]  
100%|██████████| 16/16 [00:00<00:00, 35.38it/s]  
Finished epoch 23 of 30: Train Loss = 0.783    Train Acc = 0.733    Test Acc = 0.603

100%|██████████| 78/78 [00:05<00:00, 14.80it/s]  
100%|██████████| 16/16 [00:00<00:00, 36.39it/s]

Finished epoch 24 of 30: Train Loss = 0.782    Train Acc = 0.734    Test Acc = 0.611

100%|██████████| 78/78 [00:05<00:00, 14.47it/s]

100%|██████████| 16/16 [00:00<00:00, 35.63it/s]

Finished epoch 25 of 30: Train Loss = 0.758    Train Acc = 0.738    Test Acc = 0.602

100%|██████████| 78/78 [00:05<00:00, 14.69it/s]

100%|██████████| 16/16 [00:00<00:00, 33.87it/s]

Finished epoch 26 of 30: Train Loss = 0.749    Train Acc = 0.749    Test Acc = 0.620

100%|██████████| 78/78 [00:05<00:00, 14.63it/s]

100%|██████████| 16/16 [00:00<00:00, 37.77it/s]

Finished epoch 27 of 30: Train Loss = 0.714    Train Acc = 0.759    Test Acc = 0.622

100%|██████████| 78/78 [00:05<00:00, 14.56it/s]

100%|██████████| 16/16 [00:00<00:00, 36.32it/s]

Finished epoch 28 of 30: Train Loss = 0.696    Train Acc = 0.765    Test Acc = 0.610

100%|██████████| 78/78 [00:05<00:00, 14.68it/s]

100%|██████████| 16/16 [00:00<00:00, 36.12it/s]

Finished epoch 29 of 30: Train Loss = 0.703    Train Acc = 0.760    Test Acc = 0.616

100%|██████████| 16/16 [00:00<00:00, 33.27it/s]

Test accuracy: 61.64%

